

ted2-3

June 7, 2017

1 Exercise 2: Part 3

2 Import the necessary libraries and open the data set

```
In [215]: import pandas as pd
import numpy as np
from sklearn.model_selection import KFold
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.cross_validation import train_test_split, ShuffleSplit, cross_val_score
from sklearn import preprocessing
from sklearn.naive_bayes import MultinomialNB
from sklearn import svm
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import csv
from math import log

df = pd.read_csv('train.tsv', sep='\t')
target = df["Label"]
```

3 Entropy Calculation Function

```
In [216]: def entropy(data):
    count = data.shape[0]
    if(count==0):
        return 0
    good = len(data[data["Label"]==1])
    bad = len(data[data["Label"]==2])
    goodper = good / count
    badper = bad/count
    goodentropy =0.0
    if(good>0):
        goodentropy = -(goodper*log(goodper, 2))
    badentropy =0.0
    if(bad>0):
```

```

        badentropy = -(badper*log(badper, 2))
        entropy = goodentropy+badentropy
        return entropy

```

4 Calculate the train set's category.

```

In [217]: data_set_entropy=entropy(df)
          count = len(df)

```

5 Because those 3 numerical attributes have a big variety of values we have to convert them to categorical.

```

In [218]: Attributes = list(df.columns.values)
          Attributes.remove('Id')
          Attributes.remove('Label')

          ndf = df.copy()

          NumericalAttributes = ['Attribute2', 'Attribute5', 'Attribute13']
          for Attribute in NumericalAttributes:
              ndf[Attribute] = pd.qcut(df[Attribute], 5)

```

6 Calculate the Information Gain for each Attribute.

```

In [219]: Attributes = list(df.columns.values)
          Attributes.remove('Id')
          Attributes.remove('Label')

          InformationGain_list=[]
          for Attribute in Attributes:
              Attribute_values= list(ndf[Attribute].unique())
              attr_entropy = 0.0
              for value in Attribute_values:
                  value_set= ndf[ndf[Attribute]==value]
                  attr_entropy += (len(value_set)/count)*entropy(value_set)
              InformationGain_list.append((Attribute, data_set_entropy-attr_entropy))

```

7 Sort the Information Gain List, convert the categorical to numerical.

```

In [220]: InformationGain_list.sort(key=lambda tup: tup[1])

          categories = ["Attribute1", "Attribute3", "Attribute4", "Attribute6", "Attribute7", "Attribu
          target = df["Label"]

```

```

new_df = df.copy()
for x in categories:
    converted = pd.Categorical(df[x])
    new_df[x] = converted.codes

```

8 Print the Array with the Attribute Removal Queue and the Information Gain for each one.

```

In [221]: for n in range(len(InformationGain_list)):
           print('Exclude No: ',n,' \tAttribute: ', InformationGain_list[n][0], ' \tInformation Gain: ', InformationGain_list[n][1])

```

Exclude No: 0	Attribute: Attribute18	Information Gain: 0.000129665701927850
Exclude No: 1	Attribute: Attribute11	Information Gain: 0.000220571349274112
Exclude No: 2	Attribute: Attribute19	Information Gain: 0.001202862591077602
Exclude No: 3	Attribute: Attribute16	Information Gain: 0.002395770112591733
Exclude No: 4	Attribute: Attribute17	Information Gain: 0.002940316631288131
Exclude No: 5	Attribute: Attribute10	Information Gain: 0.005674399790160045
Exclude No: 6	Attribute: Attribute14	Information Gain: 0.007041506325139002
Exclude No: 7	Attribute: Attribute8	Information Gain: 0.007330500076830004
Exclude No: 8	Attribute: Attribute20	Information Gain: 0.007704386546436126
Exclude No: 9	Attribute: Attribute15	Information Gain: 0.011618886823694607
Exclude No: 10	Attribute: Attribute13	Information Gain: 0.0117447128999153
Exclude No: 11	Attribute: Attribute9	Information Gain: 0.012746841156174304
Exclude No: 12	Attribute: Attribute7	Information Gain: 0.014547865230223445
Exclude No: 13	Attribute: Attribute12	Information Gain: 0.01490553087729540
Exclude No: 14	Attribute: Attribute5	Information Gain: 0.015294038701320956
Exclude No: 15	Attribute: Attribute6	Information Gain: 0.02219896605243432
Exclude No: 16	Attribute: Attribute4	Information Gain: 0.02689745203308369
Exclude No: 17	Attribute: Attribute2	Information Gain: 0.031782332193863394
Exclude No: 18	Attribute: Attribute3	Information Gain: 0.03788940622151615
Exclude No: 19	Attribute: Attribute1	Information Gain: 0.09382796302345509

9 Loop where we remove an attribute each time, we do 10-fold cross validation and we record the accuracy.

10 We keep the max accuracy in order to use it in testSet_Predictions.

```

In [222]: kf = KFold(n_splits=10)

```

```

cutted_df = df.copy()
exclude=["Id", "Label"]
xlist = []
ylist = []
max = 0.0
maxcount = 0

```

```

for count in range(len(InformationGain_list)):
    if(count>0):
        attr = InformationGain_list[count-1][0]
        exclude.append(attr)

my_df = new_df.copy()

excl =df.columns.difference(exclude)
new_df_to_use = new_df[excl]

xlist.append(len(InformationGain_list)-count)
RANDOM_STATE = 123

accuracy = 0.0

for x,y in kf.split(new_df_to_use):
    test = new_df_to_use.loc[y,new_df_to_use.columns]
    train = new_df_to_use.loc[x, new_df_to_use.columns]
    current_target = target[x]

    classifier= RandomForestClassifier(random_state=RANDOM_STATE)
    clf = classifier.fit(train,current_target)
    yPred = clf.predict(test)
    acc = accuracy_score(target[y], yPred)
    accuracy += acc

accuracyRF = accuracy/10
if(max < accuracyRF):
    max = accuracyRF
    maxcount =len(InformationGain_list)-count
ylist.append(accuracyRF)


```

11 Show the Accuracy Plot

```

In [223]: plt.plot(xlist, ylist, "ro-")
plt.title('Accuracy Plot')
plt.ylabel('Accuracy')
plt.xlabel('Attributes used')
plt.show()

```



12 We create the exclude list in order to get the max accuracy by removing them.

```
In [224]: exclude=["Id", "Label"]
          for count in range(maxcount):
              if(count>0):
                  attr = InformationGain_list[count-1][0]
                  exclude.append(attr)
```

13 Open test.tsv convert its attributes to numerical and cut the attributes we don't need.

```
In [225]: dfT = pd.read_csv('test.tsv', sep='\t')
          categories = ["Attribute1","Attribute3","Attribute4","Attribute6","Attribute7","Attribute8"]
          new_dfT = dfT.copy()

          IDs = dfT["Id"]
          for x in categories:
              converted = pd.Categorical(dfT[x])
              new_dfT[x] = converted.codes
```

```

excl = df.columns.difference(exclude)
new_df_to_use = new_df[excl]

dfneT = dfT.columns.difference(exclude)
new_dfT = new_dfT[dfneT]

category_dict = {1: 'Good', 2: 'Bad'}

```

14 Call the Random Forest Classifier, as he had the best accuracy.

15 Print the results in the testSet_Predictions.csv.

```

In [226]: classifier = RandomForestClassifier(random_state=RANDOM_STATE)
classifier.fit(new_df_to_use, target)
prediction = classifier.predict(new_dfT)

id = 0
with open('testSet_categories_Predictions.csv', 'w') as csvfile:
    fieldnames = ['ID', 'Label']
    writer = csv.DictWriter(csvfile, fieldnames = fieldnames, delimiter='\t')
    writer.writeheader()
    for i in range(len(prediction)):
        writer.writerow({'ID': IDs[i], 'Label': category_dict[prediction[i]]})
        id += 1
print('Created testSet_categories_Predictions.csv')

```

Created testSet_categories_Predictions.csv